

Econometrics for Political Analysis
Volume I: Causal Inference

Kevin Lingfeng Li

Table of contents

Preface	3
I Econometrics and Causal Inference	4
1 Causal Inference	5
1.1 Introduction to Econometrics	5
1.2 Potential Outcomes Framework	6
1.3 Causal Estimands	6
1.4 Estimators, Bias, and Variance	7
1.5 Naive Estimator and Selection Bias	8
2 Randomised Controlled Trials	10
2.1 Random Assignment and Experiments	10
2.2 Estimation of Causal Effects	10
2.3 Blocking and Stratified Experiments	11
2.4 Uncertainty and Standard Errors	12
2.5 Confidence Intervals and Hypothesis Testing	13
2.6 Validity and Limitations of Randomised Experiments	14
II Multiple Linear Regression	16
3 Regression Models and Estimation	17
3.1 Introduction to Regression	17
3.2 Specification of the Linear Model	18
3.3 Bivariate Regression Estimation Mechanics	19
3.4 Multivariate Regression Estimation	21
3.5 Matrix Notation of Regression	22
3.6 OLS Estimator with Linear Algebra	23
4 Interpretation and Hypothesis Testing	24
4.1 Interpretations of Coefficients	24
4.2 Model Summary Statistics	25
4.3 Uncertainty and Confidence Intervals	27
4.4 Hypothesis Testing of Parameters	27
4.5 F-Tests of Nested Models	28
4.6 Linear Regression in R	29
5 Other Explanatory Variables	32
5.1 Polynomial Transformations	32
5.2 Logarithmic Transformations	34
5.3 Binary Explanatory Variable	35
5.4 Polytomous Explanatory Variable	35

5.5	Interaction Effects	37
5.6	Explanatory Variables in R	38
III	Regression for Causal Inference	39
6	Gauss-Markov Assumptions	40
7	Panel and Clustered Data	41
7.1	Hierarchical Data	41
7.2	Fixed Effects	41
7.3	Further Approaches	41
8	Matching Regression	42
9	Weighted Regression	43
10	Partial Identification and Sensitivity Analysis	44
IV	Instrumental Variables Estimator	45
11	Instrumental Variables Framework	46
12	Other Instrumental Variable Designs	47
12.1	Continuous Independent Variables	47
12.2	Examiner Designs	47
12.3	Shift-Share Bartik Instruments	47
12.4	Recentered Instruments	47
V	Quasi-Experimental Methods	48
13	Regression Discontinuity	49
14	Simple Differences-in-Differences	50
15	Generalised Differences-in-Differences	51
16	Survey Experiments	52

Preface

This book is the first volume in a series on **Econometrics for Political Analysis**:

- [Volume I: Causal Inference](#) (this book) introduces many of the core statistical concepts and methods used in Politics research. This book mainly covers Linear Regression models and techniques of Causal Inference.
- [Volume II: Advanced Topics](#) expands on the topics covered in the first book, and covers logit and count regression models, time series models, and latent variable models.
- [Volume III: Big Data](#) introduces modern statistical learning techniques that have improved prediction accuracy. The book starts off with prediction and classification methods, then covers model validation, before ending on text mining and quantitative text analysis.

This book is designed to be both an approachable, but also rigorous, introduction to Econometrics and the use of statistical methods, specifically for the analysis of political institutions and actors. This volume specifically covers most topics you would study in a typical econometrics undergraduate course in an economics department, and is likely more in depth than what you would get in a political science department.

I assume a basic understanding of statistics, up to correlation. I also assume familiarity with calculus and some linear algebra - although this is only necessary if you are interested in how the methods work. If you are just interested in implementing methods and interpreting them, the mathematics is less important, however, I still recommend people to study the mathematics behind econometric techniques, as it often helps us understand the techniques better.

To see what mathematics is specifically required, or to refresh on the mathematics needed, consult the **Quantitative Methods** sequence (particularly [the first volume](#)).

This book will also use the R language for some calculations, so a basic understanding of the language is useful - however, this is not needed if you just want an understanding of the methods. This book will not teach the R language from scratch.

Part I

Econometrics and Causal Inference

Chapter 1

Causal Inference

1.1 Introduction to Econometrics

Econometrics is the field of applying statistical methods to analyse real-world economic and social science data. While econometrics was initially pioneered by economics, the techniques econometricians developed have been adopted by most of the social sciences, including Political Science.

Econometrics has two primary goals:

1. **Causal Inference:** Establishing how one feature directly causes another feature. This is essential to understanding the world around us and designing better policies.
2. **Predictive Inference/Forecasting:** Given data we have, how can we predict the values of data we do not have? For example, what will sales be next year? GDP? Who will win the next election? What are the likely costs/effects of a policy?

This book will cover both topics, but will focus on causal inference, which is generally considered the most important role of econometrics in the social sciences. After all, understanding the causes of things is critical to designing better policy and understanding the world around us.

An important thing to note about causal inference is that correlation does not equal causation:

- For example, ice cream sales are strongly correlated with shark attack frequency. Does that mean ice cream sales directly cause shark attacks? No - they happen to be correlated, because there is another variable, the warm weather and people going to the beach, that causes both to rise at the same time. Ice cream sales itself does not have any independent affect on shark attacks, they just happen to be correlated.

When correlation does not equal causation, that correlation is called a **spurious correlation**. [Tyler Vigen](#) has an entire website showcasing different spurious correlations if you want more examples. As social scientists, we want to avoid spurious correlations, and establish causal effects.

In the first section of this book, we will establish what a causal effect even is, and discuss how randomised experiments are the “gold standard” of establishing causal relationships. Then, we will explore why randomised experiments are often impossible in the social sciences. The rest of the book will focus on techniques that try to find causal effects without randomised experiments.

1.2 Potential Outcomes Framework

A **causal effect** is a change in some feature of the world Y , that would directly result from a change in some other feature D . Essentially, change in D causes change in Y .

This causal effect implies that there are **potential outcomes**:

- Imagine that there are 2 worlds, that are exactly the same until treatment D occurs. In one world, you get the treatment D , and the other world, you do not get this treatment.
- Since these 2 worlds are identical besides the treatment D , the difference between the world's Y outcomes are the effect of our treatment D .
- In the real world, we only observe one of these realities - either a unit i gets, or does not get, the treatment. The other world that we do not observe is called a **counterfactual**.

Causal States are these hypothetical states of the world. The control state $D = 0$ is the world where a unit does not receive the treatment D . The treatment state $D = 1$ is identical to the control state, with the only exception that a unit receives the treatment D .

For each unit of observation i , we can define two potential outcomes:

- Y_{1i} is the potential outcome for unit i , given it is in the treatment state $D_i = 1$.
- Y_{0i} is the potential outcome for unit i , given it is in the control state $D_i = 0$.

Thus, the **individual causal effect** τ of the treatment D for any unit i is $\tau_i = Y_{1i} - Y_{0i}$. This is because since the two states of the world are identical except for treatment D , the resulting difference must be as a result of treatment D .

However, in the real world, we do not have parallel worlds (unfortunately) - either unit i gets the treatment $D_i = 1$, or does not get the treatment $D_i = 0$. Thus, the **observed Y outcome** (in the real world) of any unit i is given by the equation:

$$Y_i = D_i \times Y_{1i} + (1 - D_i) \times Y_{0i}$$

This equation might be a little abstract, however, it is easy to understand by plugging D_i in:

$$\begin{aligned} [Y_i | D_i = 0] &= 0 \times Y_{1i} + (1 - 0) \times Y_{0i} = Y_{0i} \\ [Y_i | D_i = 1] &= 1 \times Y_{1i} + (1 - 1) \times Y_{0i} = Y_{1i} \end{aligned}$$

So intuitively, if the observation is in the control state $D_i = 0$, we observe the controlled state potential outcome Y_{0i} . When an observation is in the treatment state $D_i = 1$, we observe the treatment state potential outcome Y_{1i} .

1.3 Causal Estimands

The fundamental problem of causal inference is that we only can observe one of the two parallel universes at the same time. For example, if you get treatment D , we cannot observe the world

where you do not get treatment D . Thus, we cannot estimate individual effects of the causal treatment. However, there are other estimands we can use.

- Note: **Estimand** is the quantity we are trying to estimate (i.e. what we are interested in).

Since we cannot observe the individual treatment effects, we can change our estimand to the **average treatment effect (ATE)**. This is exactly what it sounds like - the treatment effect of all units averaged:

$$\tau_{ATE} = \mathbb{E}[\tau_i] = \mathbb{E}[Y_{1i} - Y_{0i}] = \frac{1}{n}(\sum Y_{1i} - \sum Y_{0i})$$

There are also other treatment effects we can use to estimate.

The **average treatment effect on the treated (ATT)** is the treatment effect of only units who received the treatment $D_i = 1$

$$\tau_{ATT} = \mathbb{E}[Y_{1i} - Y_{0i} | D_i = 1]$$

The **average treatment effect on the controls (ATC)** is the treatment effect of units who only did not receive the treatment $D_i = 0$

$$\tau_{ATC} = \mathbb{E}[Y_{1i} - Y_{0i} | D_i = 0]$$

The **conditional average treatment effect (CATE)** is the treatment effect of units, given they have some other variable X value. For example, if X is gender, the CATE could be the treatment effect on only females.

$$\tau_{CATE} = \mathbb{E}[Y_{1i} - Y_{0i} | X = x]$$

1.4 Estimators, Bias, and Variance

An estimator is the procedure by which we estimate the estimand we are trying to find.

- Remember, the above causal estimands are not directly calculable, since we do not see the counterfactual potential outcome. Thus, we have to estimate them with an estimator.

Bias is when an estimator consistently and systematically poorly estimates the estimand you are trying to get.

- Or in other words, the estimator's average estimate of our estimand (over many tries of estimation), is not actually the true value of the estimand. That means something is consistently off with our estimator - we might be consistently overestimating by 5%, or underestimating, etc.
- Or more intuitively, imagine you are trying to hit a bullseye in archery. Bias is when you might be very accurate, but aiming in the wrong place, thus not hitting the bullseye. A biased estimator is essentially that - we are consistently and systematically making a mistake when estimating the quantity in question.

Variance is the difference between our estimations derived from our estimates - i.e. how consistent our estimator is.

- For example, you might have an unbiased estimator, where our average estimate is the actual causal estimand. However, while the average is correct, the variance of our estimates is very wide.
- Or more intuitively, in the archery example, we are aiming correctly at the bullseye, however, the wind and our muscles are unpredictable, so each shot might be slightly off in different directions. If we average all our shots, we are hitting the middle, but not each individual shot is in the bullseye.

Ideally, we want an unbiased estimator that has low variance. We will explore many different types of estimators for causal effects throughout this book, each with its own bias and variance.

1.5 Naive Estimator and Selection Bias

Naive Estimator

The **naive estimator** is an estimator that only compares our observed outcomes, without any comparison to the counterfactual potential outcomes. This is often what many people initially do when trying to find a causal effect. Essentially, we are comparing units that are assigned to treatment, and the units that are not assigned to treatment, and their observed outcomes.

$$\mathbb{E}[Y_{1i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0]$$

Or more intuitively, the average observed outcome Y of those in the treatment group, minus the average observed outcome Y of those in the control group.

However, this naive estimator is a bad idea. Why?

- Remember, our treatment effects are supposed to be comparing to two potential outcomes of the same unit. We are supposed to compare Y_{1i} to Y_{0i} .
- However, in this scenario, we are not comparing the potential outcomes of the same individual. We are comparing the outcome of some observation A in treatment Y_{1A} and the outcome of some other observation B in control Y_{0B} .
- But what if observation A and B are different? Their outcomes may not be due to the treatment D , but because of the differences between A and B .
- This is why counterfactual comparison is important - when we compare the potential outcomes of the same unit in control and treatment groups, we can be confident of the affect of the treatment, since it is the same unit of observation for both groups.

We can prove this mathematically. We start with the naive estimator:

$$\mathbb{E}[Y_{1i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0]$$

Then, we do a little algebra trick - we add a new term to this equation, and then subtract the

same term. The two new terms thus cancel each other out to 0.

$$= \mathbb{E}[Y_{1i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0] + \mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 1]$$

Then, we rearrange the terms, then simplify, getting the result:

$$\begin{aligned} &= \mathbb{E}[Y_{1i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 1] + \mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0] \\ &= \mathbb{E}[Y_{1i} - Y_{0i}|D_i = 1] + \mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0] \end{aligned}$$

If we look at the final result, we can divide it into 2 parts:

1. The first part, $\mathbb{E}[Y_{1i} - Y_{0i}|D_i = 1]$, is the average treatment effect of the treated τ_{ATT} that we introduced previously.
2. The second part $\mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0]$ is what we call the **selection bias**. Intuitively, it is the difference between the treatment and control groups prior to the treatment taking place (hence the potential outcome being Y_{0i}).

Selection Bias and Confounding Variables

These differences between the treatment and control groups prior to treatment can actually explain some of our results, which is why our results with the naive estimator are **biased**.

For example, if we are measuring the question *does going to the hospital make you more healthy*, and we simply measured the outcomes of people who went to the hospital and did not go to the hospital, we might see that in general, people who did not go to the hospital are healthier!

- Does this mean that going to the hospital makes you unhealthier? No! It is because more unhealthy people choose to go to the hospital in the first place. Thus, the hospital has generally more unhealthy individuals in it. The hospital might perform miracles on these people, but they are still not as healthy as the healthy people who did not need to go to the hospital.
- The differences between the people who chose to go to the hospital versus the people who did not go to the hospital explains the differences in our outcome, not the actual treatment that the hospital provided. This is selection bias - when our treatment and control groups are fundamentally different and unequal even prior to treatment.

A **confounder** is a variable that is explaining the differences in the treatment and control groups. For example, smoking might be a confounding variable in the example above - people who smoke more often will go to the hospital, and will have worse outcomes than people who did not smoke and did not go to the hospital.

The naive estimator will capture the effect of these confounders, which we do not want - we want to isolate the effect of our treatment D . To make an accurate causal claim, we must get rid of confounding variables and selection bias. How do we do this? - Randomisation, which we will cover next.

Chapter 2

Randomised Controlled Trials

2.1 Random Assignment and Experiments

The **assignment mechanism** is how we decide which observations receive the treatment D . In the last chapter, we discussed how the Naive Estimator is biased, because of selection bias.

However, we can address this concern by randomly assigning units into either the treatment or control group. Since by random selection, each observation has an equal likelihood of being in the treatment or control, we should expect the treatment and control groups to be similar - thus eliminating selection bias.

Mathematically, we assume that after randomization, the potential outcomes are independent of treatment/control status:

$$\mathbb{E}[Y_{1i}|D_i = 1] \approx \mathbb{E}[Y_{1i}|D_i = 0] \text{ and } \mathbb{E}[Y_{0i}|D_i = 1] \approx \mathbb{E}[Y_{0i}|D_i = 0]$$

Or in other words, the potential outcomes are the same between control and treatment groups. If this assumption is met, then we can use the naive estimator to estimate the treatment effect. Remember the naive estimator:

$$\mathbb{E}[Y_{1i} - Y_{0i}|D_i = 1] + \mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0]$$

The selection bias term $\mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0] = 0$ under the above assumptions of randomisation. Thus, there is no longer selection bias, and confounding variables will have been accounted for.

2.2 Estimation of Causal Effects

After we have randomised our control and treatment groups, how do we identify the causal effect in a randomised controlled experiment?

Our observed potential outcomes in our randomised experiments are Y_{1i} and Y_{0i} . We know that the treatment group D_i does not affect our potential outcomes (as shown in the previous section).

Thus we know that:

$$\begin{aligned}\mathbb{E}[Y_{1i}|D_i = 1] &= \mathbb{E}[Y_{1i}] \\ \mathbb{E}[Y_{0i}|D_i = 0] &= \mathbb{E}[Y_{0i}]\end{aligned}$$

Now that we have $\mathbb{E}[Y_{1i}]$ and $\mathbb{E}[Y_{0i}]$, we can calculate the average treatment effect:

$$\tau_{ATE} = \mathbb{E}[\tau_i] = \mathbb{E}[Y_{1i}] - \mathbb{E}[Y_{0i}] = \bar{Y}_t - \bar{Y}_c$$

Where \bar{Y}_t is the average Y value of the treatment group, and \bar{Y}_c is the average Y value of the control group. Thus, the causal effect is simply a difference of means between the treatment and control group.

2.3 Blocking and Stratified Experiments

Experimental Design

Blocking, also called stratified experiments, is an extension of the random experiment to deal with some common issues.

Imagine that you have four units in your experiment that you have to assign to treatment/control. Their pre-treatment outcomes are $Y_{0i} = \{2, 2, 8, 8\}$. This means that you have a 1/3 chance to end up with the random assignment of $\{2, 2\}$ in one group and $\{8, 8\}$ in the other group.

- This is a major issue! After all, the core assumption of random experiments is that randomisation makes the treatment and control group similar, eliminating selection bias.

With blocking, you can prevent this from happening. Before randomisation, you separate your sample of N units into J subgroups. Then, within each group, randomly assign units to treatment and control group (essentially, smaller randomised experiments within a bigger experiment).

- For example, we could divide our prior example into 2 subgroups: $\{2, 2\}$ and $\{8, 8\}$. Then, within each group, randomly assign one observation to treatment, and one to control. Thus, we are guaranteed to get units from both subgroups in both our treatment and control groups.

Estimation of Causal Effects

To estimate our effects, we will have to take the weighted average of each subgroup's average treatment effect (ATE), with the weights being the proportion of units each group accounts for. Mathematically:

$$\tau_{ATE} = \sum_{j=1}^J \frac{N_j}{N} \tau_j$$

Where N is the total number of observations, J is the total number of subgroups, j is one of the subgroups, N_j is the number of units within subgroup j , and τ_j is the ATE of the subgroup j .

2.4 Uncertainty and Standard Errors

Intuition of Uncertainty

Remember how we randomly assigned units to treatment or control? What if we ran the experiment again? The treatment and control groups would very likely not be exactly the same, and thus, we would get a slightly different causal effect. Thus, we have some uncertainty with our causal estimate - re-running the experiment might result in a different answer.

The ATE we have calculated is only our specific sample average treatment effect (SATE), often notated $\hat{\tau}_{ATE}$ or \hat{ATE} .

- Why sample? Well, through random assignment, you are basically “randomly sampling” potential outcomes - since randomly choosing one unit to be in treatment/control means not seeing the other counterfactual potential outcome.

Thus, we need some mechanism to account for sampling variability and how rerunning the experiment might result in slightly different results. We do this with sampling distributions and standard errors.

Sampling Distributions and Standard Error

A **sampling distribution** is a hypothetical construct that is useful to understanding how many of our statistical techniques work. A sampling distribution is as follows.

- Imagine that we take a sample from a population (or some random assignment mechanism). Then, we find the average treatment effect of the sample $\hat{\tau}_{ATE}$. That is a **sample estimate**, which is often notated $\hat{\theta}$. (I use θ , since this idea of uncertainty can be applied to any estimate, not just average treatment effect).
- Then, let us take another sample from the same population (or do another random assignment), and find the sample estimate. This will be slightly different than the first sample, since we are randomly sampling. That is another sample estimate. We keep taking samples from the same population (more random assignments), and getting more and more sample estimates.
- Now, let us plot all our sample estimates $\hat{\theta}$ (different $\hat{\tau}_{ATE}$ values) into a “histogram” or density plot. The x axis labels the possible $\hat{\tau}_{ATE}$ values, and the y axis is how frequently a specific sample estimate occurs. We will get a distribution, just like a random variable distribution.
- That distribution is the **sampling distribution**

A Sampling distribution is the imaginary distribution of estimates, if we repeated the sampling and estimation process many, many times.

The **standard error** is the standard deviation of the sampling distribution. It is often notated $SE(\hat{\theta})$. The computer/software we use will calculate this for us.

2.5 Confidence Intervals and Hypothesis Testing

Confidence Intervals

Since there is variability of estimates between samples, we have to create an interval around our sample estimate $\hat{\theta}_j$ to account for this uncertainty. We assume our estimated $\hat{\theta}_j$ is the centre of this distribution, then add some “buffer” to both sides. The confidence interval’s lower and upper bounds are defined as, given a confidence level of 95% (the standard confidence level):

$$\hat{\theta}_j \pm 1.96 \times \widehat{se}(\hat{\theta}_j)$$

- $\widehat{se}(\hat{\theta}_j)$ is the standard error of our estimate of how precisely we have estimated the true value of θ_j , introduced in the previous section.
- Why 1.96? It is because in a normal distribution, 95% of the data is contained within 1.96 standard deviations, and Central Limit Theorem states that sampling distributions are normally distributed. (The 1.96 can vary based on sample size however, if we do not meet Central Limit Theorem criteria, but that is not important for now).

Confidence intervals say that if we repeated the sampling and estimation process many many times (like we did for our sampling distribution), 95% of the confidence intervals we construct from our samples, would correctly contain the true θ_j .

Every value in a given confidence interval is a plausible value of the true θ_j in the population. The most important thing is if 0 is included within the confidence interval. $\theta = 0$ means that there is no causal effect.

Hypothesis Testing

In academia, we are conservative - this means we do not claim we have found a new theory, unless we are quite confident that the old theory was not true. The old theory is called our **null hypothesis**, often notated H_0 . This is the old theory that we are trying to disprove.

- Most often, the “old theory” we are trying to disprove is that *there is no causal effect of D on Y* (since it is very rare we are studying something that has already been proven). No relationship means that $\theta_j = 0$

The new theory we have come up with, and are trying to prove, is called the **alternate hypothesis**, often notated H_1 or H_a .

- In general, our new hypothesis is that *there is a causal effect of D on Y*, or $\theta_j \neq 0$

We assume that the null hypothesis is true, unless we are 95% confident that we can reject the null hypothesis, and only then, can we accept the alternative hypothesis (the new theory we proposed).

How do we actually test these hypotheses? First, we have to calculate a t-test statistic. The formula for such a statistic is the parameter divided by its standard error (calculated by the computer):


$$t = \frac{\hat{\theta}_j}{\widehat{se}(\hat{\theta}_j)}$$

The t-test statistic basically tells us how far the parameter we tested is from 0, in terms of standard errors of the parameter.

Then, we have to consult a t-distribution (see section 2.4). T-distributions only has one parameter - degrees of freedom, which is calculated by the number of observations n , minus the number of variables k , then minus 1: $DF = n - k - 1$

With the degrees of freedom, we can find the corresponding t-distribution, labeled t_{n-k-1} . Then, we start from the middle of that t distribution, and go the *number of standard errors* away based on the t-test statistic. We do this on both sides from the middle of the t-distribution.

Once we have found that point, we find the probability (area under the distribution) of a t-test statistic of ours, or more extreme, could occur. The figure below, with its blue highlighted area, shows this probability:



The area highlighted is our p-value. Essentially, a p-value is how likely we are to get a test statistic at or more extreme than the one we got for our estimated θ_j , given the null hypothesis is true.

- So if the p-value is very high, there is a high chance that the null hypothesis is true.
- If the p-value is very low, then there is a low chance that the null hypothesis is true

Generally, in the social sciences, if the p-value is less than 0.05 (5%), we can **reject the null hypothesis**, and conclude the alternate hypothesis.

2.6 Validity and Limitations of Randomised Experiments

There are two types of validity in Randomised Experiments: Internal and External validity.

Internal validity is about if our experiment accurately captures the average causal effect of our units in our experiment. Essentially - did we estimate the causal effect correctly? Some things that can cause lack of internal validity include:

1. Failure of randomisation: if treatment and control groups are not similar, we violate assumptions of random experiments and that will include selection bias in our estimates.
2. Non-Compliance: Sometimes, our subjects that are assigned to treatment, refuse to comply with the treatment (we cannot force them to). This will mess up the average treatment effect, since some units did not properly undergo the treatment.
3. Attrition: Sometimes, outcomes cannot be measured for some study participants, for example, if they drop out or refuse to answer. This is concerning - because the people who drop out might have some common characteristic (confounding variable), and we will miss this entirely in our estimation.

External Validity is about the generalisation of our conclusions - we know the effect on our experimental subjects, but does this causal effect apply to other units across the world?

- For example, if we do a study in Japan, can we assume that the same effects are applicable in the US? South Africa?
- Generally to obtain this, you want the units included in your observation to be representative of the larger units you want to apply your results to. For example, if you are measuring the causal effect of some treatment on Americans, you want your subjects to be representative of Americans as a whole.

Finally, Randomised Experiments have some **limitations**.

1. Ethical limitations: sometimes, it is unethical to have units take potentially dangerous treatments, or have some units not undergo potential benefits of treatment. We are essentially randomly selecting what happens to people's lives.
2. Practical limitations: often, running experiments is just not possible. For example, let us say you want to see if democracy increases economic growth. To do this, you would need to randomly assign countries to democracy or autocracy (control) groups. But let us be honest, you can't force Canada to be a dictatorship against their will. Often, we will have to use **observational studies** - where we do not control assignment of treatment.

The rest of this course focuses on observational studies. We will introduce regression - one of the most powerful tools in causal inference. Then, we will introduce extensions to regressions that help us address some of the limitations regressions have.

Part II

Multiple Linear Regression

Chapter 3

Regression Models and Estimation

Note: This chapter's later sections may become quite complex. However, do not worry if you do not understand section 3.3 and beyond, as only the sections 3.1-3.2 are absolutely essential.

3.1 Introduction to Regression

As we discussed in the previous sections, a randomised experiment is the best way of establishing causal relationships. This is because randomise treatments can get rid of the effect of confounding variables. However, in the social sciences, randomisation is often not possible.

Luckily, while we often cannot randomise treatment, we can account for the affect of confounding variables with a new tool: linear regression.

Linear regression allows us to estimate a model with both our explanatory and outcome variables, as well as a series of **control** variables. By including confounding variables as control variables in our regression model, we can (in theory), isolate the effect of our explanatory variable on our outcome variable.

- In reality, as we will discuss in the later parts of this book, there are often thousands of control variables, many that are not possible to control for. We will introduce extensions to the linear model to deal with these.

Regression is also a powerful tool for predictive inference. In fact, it forms much of the building blocks of more complicated data science and machine learning methods.

Before we dive into the linear model, here are some conventional notation that is important:

- The **response variable** (dependent variable) is notated Y . In this book, we will only have one response variable.
- The **explanatory variable** (independent variable) is notated X . There is often more than one explanatory variable, so we denote them with subscripts X_1, X_2, \dots, X_k . We sometimes also denote all explanatory variables as the vector \vec{X}
- Note: our treatment variable (for causal inference) D is considered one of the explanatory variables \vec{X} . We typically define the first of these explanatory variables X_1 as the treatment variable, and all others X_2, \dots, X_k as control variables.

3.2 Specification of the Linear Model

A regression model is the specification of the conditional distribution of Y , given \vec{X} . Essentially, it is stating that the distribution of possible Y outcomes depends on the value of \vec{X} . The linear regression model focuses on the **expected value** of the conditional distribution of Y given \vec{X} .

- I say **distribution** because there are often a range of Y outcomes, each with their own probabilities, for any given X . For example, if X was age and Y was income, at age $X = 30$, not every single 30 year old makes the same amount of money. There is some distribution of incomes Y at age $X = 30$. The regression focuses on the expected value of Y at some X .

Form of the Linear Model

Suppose we have a set of observed data, with response variable Y , and a number of X variables for n number of observations. Thus, we will have n number of pairs of (X_i, Y_i) observations. The linear model takes the following form:

$$E[Y_i|\vec{X}_i] = \alpha + \beta_1 X_{1i} + \dots + \beta_k X_{ki}$$

- Where $E[Y_i|\vec{X}_i]$ is the expected value of the conditional distribution $Y_i|\vec{X}_i$
- The distribution of $Y_i|\vec{X}_i$ has a variance $Var(Y_i|\vec{X}_i) = \sigma^2$.
- The parameters of the model are denoted by the vector $\vec{\beta}$, and contain $\alpha, \beta_1, \dots, \beta_k$

We can also write the linear model for the value of any point Y_i in our data:

$$Y_i = \alpha + \beta_1 X_{1i} + \dots + \beta_k X_{ki} + \epsilon_i$$

- Where ϵ_i is the error term function - that determines the error for each point. We will go into detail on this later.
- A key assumption (that we will discuss later) is that the error function overall is normally distributed: $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$
- Essentially, ϵ_i is another way to think about the conditional distribution of $Y_i|\vec{X}_i$, and how not every 30 year old makes the exact same income - there is some variation (and error).

Parameters and Estimation

In our model, we have parameters $\alpha, \beta_1, \dots, \beta_k$ that need to be estimated in order to create a best-fit line we can actually use. We estimate the parameters and fit the model by using our observed data points (Y_i, \vec{X}_i) , and fitting a best fit line to these points. Our result should take the following form:

$$\hat{Y}_i = \hat{\alpha} + \hat{\beta}_1 X_{1i} + \dots + \hat{\beta}_k X_{ki}$$

- Where \hat{Y}_i is our prediction of the value of Y , given any set of \vec{X} values.
- Notice the error term ϵ_i is not present. This is because of our prior assumption that the error term $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$, which says the expected value of ϵ_i is $E[\epsilon_i] = 0$.

However, how do we determine the estimates of our parameters $\alpha, \beta_1, \dots, \beta_k$? **The Ordinary Least Squares Estimator (OLS)**. OLS estimation attempts to **minimise the sum of squared errors** of our predicted line to our actual observed data.

What are the sum of squared errors? They are exactly what they sound like. First, calculate the difference between our estimated Y from our model \hat{Y}_i , and the actual Y_i value. Then, square that. Then sum all of those squared errors for every observation.

$$SSE = \sum (Y_i - \hat{Y}_i)^2 = \sum (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{1i} - \dots - \hat{\beta}_k X_{ki})^2$$

The final equation is obtained by substituting our fitted model in for \hat{Y}_i .

3.3 Bivariate Regression Estimation Mechanics

A bivariate regression (also called a simple linear regression) is a linear regression model with only one explanatory variable X . We noted how the **ordinary least squares** (OLS) estimator is concerned with the sum of squared errors. Let us define the sum of squared errors as a function S .

$$S(\hat{\alpha}, \hat{\beta}) = \sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta} X_i)^2$$

The OLS estimator wants to find the parameters that **minimise the sum of squared errors**. We can mathematically express this.

$$(\hat{\alpha}, \hat{\beta}) = \arg \min_{\hat{\alpha}, \hat{\beta}} \sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta} X_i)^2 = \arg \min_{\hat{\alpha}, \hat{\beta}} S(\hat{\alpha}, \hat{\beta})$$

Parameter $\hat{\alpha}$

Let us first look at the parameter $\hat{\alpha}$. How do we find what value $\hat{\alpha}$ minimises the sum of squared errors? We know through calculus, that deriving the function to find the first order condition can accomplish this:

$$\frac{\partial S(\hat{\alpha}, \hat{\beta})}{\partial \hat{\alpha}} = \frac{\partial S(\hat{\alpha}, \hat{\beta})}{\partial \hat{\alpha}} \left[\sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta} X_i)^2 \right]$$

First, ignore the summation. The partial derivative of the internal section, using chain rule, is the following:

$$\frac{\partial}{\partial \hat{\alpha}} [(Y_i - \hat{\alpha} - \hat{\beta} X_i)^2] = -2(Y_i - \hat{\alpha} - \hat{\beta} X_i)$$

But how do we deal with the summation? We know that there is the sum rule of derivatives $[f(x) + g(x)]' = f'(x) + g'(x)$. Thus, we know we just sum up the derivatives:

$$\begin{aligned} \frac{\partial S(\hat{\alpha}, \hat{\beta})}{\partial \hat{\alpha}} &= \sum_{i=1}^n [-2(Y_i - \hat{\alpha} - \hat{\beta} X_i)] \\ &= -2 \sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta} X_i) \end{aligned}$$

To find the value of $\hat{\alpha}$ that creates the minimum value of the SSE, we set the first order derivative equal to 0. We can ignore the -2, since if the sum is equal to 0, then the -2 will have no effect. Now, using properties of summation, isolate $\hat{\alpha}$ as follows:

$$\begin{aligned}
\sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta}X_i) &= 0 \\
\sum_{i=1}^n Y_i - n\hat{\alpha} - \hat{\beta} \sum_{i=1}^n X_i &= 0 \\
-n\hat{\alpha} &= -\sum_{i=1}^n Y_i + \hat{\beta} \sum_{i=1}^n X_i \\
\hat{\alpha} &= -\frac{1}{n} \left[-\sum_{i=1}^n Y_i + \hat{\beta} \sum_{i=1}^n X_i \right] \\
\hat{\alpha} &= \frac{1}{n} \sum_{i=1}^n Y_i - \hat{\beta} \sum_{i=1}^n X_i \\
\hat{\alpha} &= \bar{Y} - \hat{\beta}\bar{X}
\end{aligned}$$

The final step converting to \bar{Y} and \bar{X} is because of the mathematical definition of average. We will plug this $\hat{\alpha}$ equation into our solution for $\hat{\beta}$ to solve that. Once we solve $\hat{\beta}$, we will come back and calculate $\hat{\alpha}$'s value.

Parameter $\hat{\beta}$

Now, let us find the minimum $\hat{\beta}$ value by taking the partial derivative of the SSE function S in respect to $\hat{\beta}$ and setting it equal to 0. This is almost the same as before - use chain rule, then use sum rule to get the derivative:

$$\begin{aligned}
\frac{\partial S(\hat{\alpha}, \hat{\beta})}{\partial \hat{\beta}} &= \sum_{i=1}^n [-2X_i(Y_i - \hat{\alpha} - \hat{\beta}X_i)] \\
&= -2 \sum_{i=1}^n X_i(Y_i - \hat{\alpha} - \hat{\beta}X_i)
\end{aligned}$$

Now, let us plug in our previously solved $\hat{\alpha} = \bar{Y} - \hat{\beta}\bar{X}$, and we get:

$$\frac{\partial S(\hat{\alpha}, \hat{\beta})}{\partial \hat{\beta}} = -2 \sum_{i=1}^n [X_i(Y_i - [\bar{Y} - \hat{\beta}\bar{X}] - \hat{\beta}X_i)]$$

Once again, the -2 does not matter as same reason as before. Set equal to 0 to solve for the value of $\hat{\beta}$ that minimises SSE:

$$\begin{aligned}
0 &= \sum_{i=1}^n [X_i(Y_i - [\bar{Y} - \hat{\beta}\bar{X}] - \hat{\beta}X_i)] \\
0 &= \sum_{i=1}^n [X_i(Y_i - \bar{Y} - \hat{\beta}(X_i - \bar{X}))] \\
0 &= \sum_{i=1}^n [X_i(Y_i - \bar{Y}) - X_i\hat{\beta}(X_i - \bar{X})]
\end{aligned}$$

$$0 = \sum_{i=1}^n X_i(Y_i - \bar{Y}) - \hat{\beta}_1 \sum_{i=1}^n X_i(X_i - \bar{X})$$

Here are a few properties on summation that will help us solve this equation:

$$\begin{aligned} \sum_{i=1}^n (X_i - \bar{X}) &= 0 \\ \sum_{i=1}^n X_i(Y_i - \bar{Y}) &= \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y}) \\ \sum_{i=1}^n X_i(X_i - \bar{X}) &= \sum_{i=1}^n (X_i - \bar{X})^2 \end{aligned}$$

With these rules, we can transform what we had before into:

$$0 = \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y}) - \hat{\beta} \sum_{i=1}^n (X_i - \bar{X})^2$$

Then solve for $\hat{\beta}$ to get:

$$\hat{\beta} = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{Cov(X, Y)}{Var(X)} = \frac{\sigma_{XY}}{\sigma_x^2}$$

Thus, with this solution, we have estimated $\hat{\beta}$. If we recall, $\hat{\alpha} = \bar{Y} - \hat{\beta}\bar{X}$, so we can quickly solve for that as well. Thus, we now have $\hat{\beta}, \hat{\alpha}$, completing our estimation. We can now put everything into the fitted model:

$$\hat{Y}_i = \hat{\alpha} + \hat{\beta}X_i$$

The only thing that must be true to estimate bivariate linear regression is that $Var(X) \neq 0$.

3.4 Multivariate Regression Estimation

Multivariate regression, as introduced previously, allows us to add additional control variables:

$$Y_i = \hat{\alpha} + \hat{\beta}_1 X_{1i} + \hat{\beta}_2 X_{2i} + \dots + \hat{\beta}_k X_{ki}$$

Similar to our bivariate regression (but with additional variables), our minimisation condition is:

$$\begin{aligned} (\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2, \dots) &= \arg \min_{(\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2, \dots)} (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{1i} - \hat{\beta}_2 X_{2i} \dots)^2 \\ &= \arg \min_{(\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2, \dots)} S(\hat{\alpha}, \hat{\beta}_1, \hat{\beta}_2, \dots) \end{aligned}$$

However, instead of a best-fit line in a 2-dimensional setting, we now have a best-fit plane in a 3-dimensional space (or higher depending on the number of explanatory variables).

Taking the partial derivatives of each parameter as before, and setting them equal to 0, we get these three first order conditions:

$$\begin{aligned}
-2 \sum_{i=1}^n (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{1i} - \hat{\beta}_2 X_{2i} \dots) &= 0 \\
-2 \sum_{i=1}^n X_{1i} (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{1i} - \hat{\beta}_2 X_{2i} \dots) &= 0 \\
-2 \sum_{i=1}^n X_{2i} (Y_i - \hat{\alpha} - \hat{\beta}_1 X_{1i} - \hat{\beta}_2 X_{2i} \dots) &= 0 \\
\text{and so on for } X_{3i}, \dots, X_{ki}
\end{aligned}$$

Just like before, we can ignore the -2 in the conditions.

The 1st equation, solving for $\hat{\alpha}$, is quite simple, just as in bivariate regression. We get:

$$\hat{\alpha} = \bar{Y} - \hat{\beta}_1 \bar{X}_1 - \hat{\beta}_2 \bar{X}_2 \dots$$

The other conditions are difficult to solve by hand, so the computer will do this for us. However, it is still useful to understand what the computer is trying to do - minimise square errors.

The only properties that are required to compute the multivariate regression OLS estimates are:

1. All explanatory variables have sample variability $Var(X_j \in \bar{X}) \neq 0, \forall X_j \in X$.
2. No two explanatory variable have perfect collinearity - which means a correlation coefficient of -1 or 1, which only occurs if the two variables have the same exact values. Mathematically, $X_{ai} \neq X_{bi}, \forall i, \forall X_a, X_b \in \bar{X}$.

3.5 Matrix Notation of Regression

By expressing models in linear algebra, it becomes easier to derive formal results without requiring messy algebra and the summation operator. It is also natural, since our data is stored in spreadsheets (big matrices).

We start with the linear model:

$$Y_i = \alpha + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_k X_{ki} + \epsilon_i$$

The i th observation can be written in vector form as following:

$$y = X_i' \beta + \epsilon_i, \text{ where } \beta = \begin{bmatrix} \alpha \\ \beta_1 \\ \vdots \\ \beta_k \end{bmatrix} \text{ and } X_i = \begin{bmatrix} 1 \\ X_{1i} \\ \vdots \\ X_{ki} \end{bmatrix}$$

- The X_i' in the equation is the transpose of X_i , to make matrix multiplication possible.
- The first element of the X_i matrix is 1, since $1 \times \alpha$ gives us the first parameter in the linear model.

Since our model has n different observations of i , we can express this into vector form, with the X'_i and β being vectors within a vector.

$$\begin{pmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{pmatrix} = \begin{pmatrix} X'_1\beta + \epsilon_1 \\ X'_2\beta + \epsilon_2 \\ \vdots \\ X'_n\beta + \epsilon_n \end{pmatrix} = \begin{pmatrix} X'_1\beta \\ X'_2\beta \\ \vdots \\ X'_n\beta \end{pmatrix} + \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}$$

Since β vector appears as a common factor for all observations $i = 1, \dots, n$, we can factor it out and have an equation:

$$\begin{pmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{pmatrix} = \begin{pmatrix} X'_1 \\ X'_2 \\ \vdots \\ X'_n \end{pmatrix} \beta + \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}$$

We can expand the X'_1, \dots, X'_n vector into a matrix. Remember that each X'_1, \dots, X'_n is already a vector of different explanatory variables. Thus, we have a model:

$$Y = X\beta + \epsilon, \text{ where } X = \begin{bmatrix} 1 & X_{21} & \dots & x_{k1} \\ 1 & X_{22} & \dots & x_{k2} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & X_{2n} & \dots & X_{kn} \end{bmatrix}$$

- Where the notation for elements of X is X_{ki} , with i being the unit of observation $i = 1, \dots, n$, and k being the explanatory variables index.
- Where Y and ϵ are $n \times 1$ vectors (as seen above), and β is a $k \times 1$ vector.
- The first row of X is a vector of 1, which exists because these 1's are multiplied with α in our model.

3.6 OLS Estimator with Linear Algebra

Chapter 4

Interpretation and Hypothesis Testing

4.1 Interpretations of Coefficients

Simple Linear Regression

Simple Linear Regression is a special case of the linear model, when there is only one explanatory variable X . How do we interpret parameters $\hat{\alpha}$ and $\hat{\beta}$ that we have calculated?

$\hat{\beta}$ is the slope of the the linear model. $\hat{\beta}$ is the expected change in Y , given a one-unit increase in X .

- A positive $\hat{\beta}$ means a positive relationship, a negative $\hat{\beta}$ means a negative relationship, and $\hat{\beta} = 0$ means no relationship.
- This is only for continuous X explanatory variables. See Chapter 5 for categorical/binary explanatory variables.

$\hat{\alpha}$ is the y-intercept of the linear model. $\hat{\alpha}$ is the expected value of \hat{Y} , given $X = 0$.

Multiple Linear Regression

Multiple linear regression is when there are multiple explanatory variables X_1, X_2, \dots, X_k . How do we interpret these parameters $\hat{\alpha}$ and $\hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_k$ that we have calculated?

- Note: I will define $\hat{\beta}_j$ as any one of $\hat{\beta}_1, \dots, \hat{\beta}_k$ (the interpretation is all the same). In the model, $\hat{\beta}_j$ will be multiplied to variable X_j .

Formally, any coefficient $\hat{\beta}_j$ is the expected change in Y , corresponding to a one unit increase in X_j , holding all other explanatory variables X_1, \dots, X_k that are not X_j constant.

- Essentially, we do the same interpretation as the single linear regression, but adding the phrase “**holding all other explanatory variables constant**”.

$\hat{\alpha}$ is the expected value of \hat{Y} , given all explanatory variables X_1, \dots, X_k equal 0.

Interpreting in Terms of Standard Deviation

Sometimes, it is hard to understand what changes in Y and X mean in terms of units. For example, if democracy is measured on a 100 point scale, what does a 5 point change in democracy mean? Is it a big change, or a small change?

We can add more relevant detail by expressing the change of Y and X in standard deviations.

So, instead of the expected change of Y given one unit increase of X , we instead do the expected standard deviation change of Y , given a one standard deviation increase in X .

How do we calculate this? There is a formula! Y changes by $(SD_{X_j} \times \hat{\beta}_j)/SD_Y$, where SD represents standard deviation, and X_j is the variable whose coefficient we are interpreting.

Binary Y Variable Interpretation

If our response variable is binary, (i.e. Y only has two values, 0 and 1), then our interpretation differs slightly. We treat Y as a variable of two categories, $Y = 0$ and $Y = 1$.

$\hat{\beta}_j$ is the expected change in the probability of getting category $Y = 1$, given a one-unit increase in X_j . We could multiply this by 100 to get the expected percentage point change.

$\hat{\alpha}$ is the expected probability of getting category $Y = 1$, when $X_j = 0$.

An important note is that, in theory, for a linear regression model, Y should be continuous, not binary. In theory, we should be using logistic regression instead of linear regression (which is covered in the Advanced Statistical Methods book).

- However, in practice, researchers often do use linear regression for binary Y , simply because linear regression is easier to interpret, and easier to incorporate into more causal inference methods.

4.2 Model Summary Statistics

Estimated Residual Standard Deviation

We can derive the estimate of the **residual variance** σ^2 with this formula:


$$\hat{\sigma}^2 = \frac{\sum (Y_i - \hat{Y}_i)^2}{n - k - 1}$$

But what is the residual variance? Recall the way we write our regression model:

$$Y_i = \alpha + \beta_1 X_{1i} + \dots + \beta_k X_{ki} + \epsilon_i$$

We know that $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$. Our estimate of the residual variance $\hat{\sigma}^2$ is our estimate of the variance of the error term ϵ_i 's variance. More intuitively, it explains how spread out observed values of Y are from our prediction value $\hat{Y} = E(Y|X)$.

The figure below better showcases this in 2 different models. The red lines are our predicted regression line, and the green lines represent the distribution of our error term ϵ_i :



The residual standard deviation $\hat{\sigma}$ (square root of variance) is consistent throughout a model. This is one of the assumptions of the linear regression model - that errors are consistently distributed, no matter the value of X . This assumption is called **homoscedasticity**.

If $\hat{\sigma}$ varies depending on the value of X , then that is called **heteroscedasticity**. When this occurs, it is often a suggestion that our relationship may not be linear - and we perhaps need to try a few transformations. We will get into transformations in a later chapter.

Total Sum of Squares

The total sum of squares is the total amount of sample variation in Y :

$$\begin{aligned} TSS &= \sum (Y_i - \bar{Y})^2 \\ &= \sum (\hat{Y}_i - \bar{Y})^2 + \sum (Y_i - \hat{Y}_i)^2 \end{aligned}$$

Where TSS is the total sum of squares, SSM $\sum (\hat{Y}_i - \bar{Y})^2$ is the model sum of squares, and SSE $\sum (Y_i - \hat{Y}_i)^2$ is the sum of squared errors (that we used to fit the model).

SSM (model sum of squares) represents the part of the variation of Y that is explained by the model, while SSE (sum of squared errors) represents the part of the variation of Y that is not explained by the model (hence, why it is called error).

R-Squared Statistic

R-squared R^2 is a measure of the percentage of variation in Y , that is explained by our model (with our chosen explanatory variables). The percentage of variation in Y explained by our model would be:

$$R^2 = \frac{SSM}{TSS} = \frac{\sum (\hat{Y}_i - \bar{Y})^2}{\sum (Y_i - \bar{Y})^2}$$

Since R^2 shows how much of the variation in Y our model explains, it is often used as a metric for how good our model is - however, don't overly focus on R^2 , it is just one metric with its benefits and drawbacks.

4.3 Uncertainty and Confidence Intervals

Samples

A **sample** is a subset of a population, which ideally, can tell us something about the population. For example, we could use a sample to estimate the relationship between two features X_j and Y with a coefficient β_j . If our sample can reflect the population, then we can use the sample to learn about the true population β_j .

- In causal inference, as we have discussed before, we are sampling from the population of potential outcomes.

The gold standard of sampling procedure is a **random sample** - where individuals in the sample are selected at random from the population. In a random sample, every possible individual has an equal chance of being selected, and thus, the resulting sample is likely to be reflective of the population.

Confidence Intervals

We previously discussed confidence intervals in section 2.4. The mechanics are practically the same, but we replace the sample average treatment effect with $\hat{\beta}_j$ as our sample estimate.

Our OLS estimation has produced a $\hat{\beta}_j$ based on the data in one sample. However, remember, the sample is just a fraction of the population - and the true β_j is unlikely to be exactly the one in our sample, due to sampling variation. Thus, we have to create an interval around our estimate $\hat{\beta}_j$ to account for this uncertainty. We assume our estimated $\hat{\beta}_j$ is the centre of this distribution, then add some “buffer” to both sides:

$$\hat{\beta}_j \pm 1.96 \times \widehat{se}(\hat{\beta}_j)$$

4.4 Hypothesis Testing of Parameters

We previously discussed hypothesis testing in section 2.5. The mechanics are practically the same, but we replace the sample average treatment effect with $\hat{\beta}_j$ as our sample estimate. Generally, for regressions, our hypotheses that we test are:

- $H_0 : \beta_j = 0$ - i.e. there is no relationship between X_j and Y
- $H_1 : \beta_j \neq 0$ - i.e. there is a relationship between X_j and Y

How do we actually test these hypotheses? First, we have to calculate a t-test statistic. The formula for such a statistic is the parameter divided by its standard error (calculated by the computer):

$$t = \frac{\hat{\beta}_j}{\widehat{se}(\hat{\beta}_j)}$$

The t-test statistic basically tells us how far the parameter we tested is from 0, in terms of standard errors of the parameter.

Then, we have to consult a t-distribution (see section 2.4). We find the probability (area under the distribution) of a t-test statistic of ours, or more extreme, could occur.

Essentially, a p-value is how likely we are to get a test statistic at or more extreme than the one we got for our estimated β_j , given the null hypothesis is true.

- So if the p-value is very high, there is a high chance that the null hypothesis is true.
- If the p-value is very low, then there is a low chance that the null hypothesis is true

Generally, in the social sciences, if the p-value is less than 0.05 (5%), we can **reject the null hypothesis**, and conclude the alternate hypothesis.

4.5 F-Tests of Nested Models

An F-test is a variation on the coefficient significance tests. The standard F-test is quite simple - it tests for the significance of a model across multiple coefficients:

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_k = 0$$

$$H_1 : \text{at least one of } \beta_1, \beta_2, \dots, \beta_k \neq 0$$

Thus, we assume that all explanatory variables in the model have no relationship to the outcome variable, unless we can reject the null hypothesis and show that at least one coefficient is statistically significant.

For a standard Linear Regression, if we have any significant individual coefficients, then the F-test will also become significant. So what is the point of F-tests? Well, for some models that we will see later, such as categorical explanatory variables, the coefficient's significance levels mean something else. Thus, to see if the model is statistically significant, we need to use the F-test.

The **F-test of Nested Models** is an extension of the standard F-test, that allows us to compare different regression models. We use a smaller model as our null hypothesis, and a larger model (containing the smaller model) as our alternative hypothesis. More mathematically:

$$M_0 : E[Y] = \alpha + \beta_1 X_1 + \dots + \beta_g X_g$$

$$M_a : E[Y] = \alpha + \beta_1 X_1 + \dots + \beta_g X_g + \beta_{g+1} X_{g+1} + \dots + \beta_k X_k$$

Importantly, note how the null hypothesis model is entirely contained within the alternate hypothesis model. This must be the case - all explanatory variables in model M_0 must also be in M_a , along with additional explanatory variables in M_a .

The F-test uses the F-test statistic. This statistic compared the R^2 values of the two models. Let us say the R^2 value of M_0 is notated R_0^2 , and the R^2 value of M_a is notated as R_a^2 . The F-test statistic essentially measures the difference $R_a^2 - R_0^2$. If the difference is sufficiently large, that means the M_a model has significantly more explanatory power than M_0 .

Let SSE_a and R_a^2 denote the sum of squared errors and R^2 values of model M_a , and SSE_0 and R_0^2 for model M_0 . The total number of coefficients of M_a are k_a , and for M_0 is k_0 . Mathematically, the F-test statistic is as follows:

$$F = \frac{(SSE_0 - SSE_a)/(k_a - k_0)}{SSE_a/[n - (k_a + 1)]}$$

$$F = \frac{\frac{R_a^2 - R_0^2}{[n - (k_0 + 1)] - [n - (k_a + 1)]}}{(1 - R_a^2)/[n - (k_a + 1)]}$$

$$F = \frac{R_{\text{change}}^2/df_{\text{change}}}{(1 - R_a^2)/[n - (k_a + 1)]}$$

The sampling distribution of the F-statistic is the F distribution with parameters $k - a - k_0$ and $n - (k_a + 1)$ degrees of freedom. We then obtain the p-value from this distribution.

The p-values of the F-statistic show the following:

- If the p-value is very small, that means R_a^2 is significantly larger than R_0^2 . This is evidence against model M_0 , and in favour of the larger model M_a
- If the p-value is large, that means R_a^2 is not much larger than R_0^2 . This means there is no evidence against M_0 , and M_a is not the statistically significantly better model.

F-tests of nested models can help us determine if we should include certain extra explanatory variables. If the addition of the explanatory variables does not statistically significantly improve the performance of the model, there is little reason to include them, unless we have some other theoretical reason to include them.

4.6 Linear Regression in R

Setup

As for all of the examples in this book, we will use the “tidyverse” package. Let us also load our dataset that we will be using.

```
# if you haven't installed tidyverse, do: install.packages('tidyverse')
library(tidyverse)
democracy_data <- read_csv("data/democracy.csv")
```

Syntax of Regression

We use the `lm()` function to run a regression: The general syntax is as follows:

- Replace *model_name* with your model name, *Y* with the name of your response variable, *X1*, *X2*... with the name of your explanatory variable, and *mydata* with the name of your dataset.
- Add additional explanatory variables with more + signs, and you can remove them down to a minimum of one *X*

```
model_name <- lm(Y ~ X1 + X2 + X3, data = mydata)
summary(model_name)
```

Example of Regression

For example, let us run a model predicting *polity_2* with 2 explanatory variables

```
model2 <- lm(polity_2 ~ GDP_Per_Cap_Haber_Men_2 + Total_Oil_Income_PC,
             data = democracy_data)
summary(model2)
```

Call:

```
lm(formula = polity_2 ~ GDP_Per_Cap_Haber_Men_2 + Total_Oil_Income_PC,
    data = democracy_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-51.604	-5.790	-0.162	6.246	40.458

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.775e+00	8.554e-02	-20.75	<2e-16 ***
GDP_Per_Cap_Haber_Men_2	4.764e-04	1.084e-05	43.93	<2e-16 ***
Total_Oil_Income_PC	-1.103e-03	2.850e-05	-38.69	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.639 on 10267 degrees of freedom

Multiple R-squared: 0.1729, Adjusted R-squared: 0.1727

F-statistic: 1073 on 2 and 10267 DF, p-value: < 2.2e-16

We can see the output. In the coefficients table:

- The *Intercept - Estimate* is $\hat{\alpha}$
- The *GDP_Per_Cap_Haber_Men_2 - Estimate* is the coefficient $\hat{\beta}_1$.
- The *Total_Oil_Income_PC - Estimate* is the coefficient $\hat{\beta}_2$.

Now, look further right in the coefficients table:

- Under the *Std. Error* column is the respective standard errors for both β coefficients. This is used to calculate the t-statistic
- Under the *t value* column is the t-test statistics for both β coefficients.
- Under the *Pr(>|t|)* column is the p-values for each test statistic. Stars symbolise statistical significance.

Underneath, the *Residual Standard Error* is the Residual Standard Deviation $\hat{\sigma}^2$, and the *Multiple R-Squared* is the R^2 value. Finally, at the very bottom, we have the p-value of the F-test if $H_0 : \vec{\beta} = 0$.

Confidence Intervals in R

To calculate confidence intervals, we can use the `confint()` command, and simply input the name of our model within:

```
confint(model2)
```

	2.5 %	97.5 %
(Intercept)	-1.9422959405	-1.6069289543
GDP_Per_Cap_Haber_Men_2	0.0004551738	0.0004976894
Total_Oil_Income_PC	-0.0011586240	-0.0010468834

Here we can see that R has generated confidence intervals for $\hat{\alpha}_1$, $\hat{\beta}_1$, and $\hat{\beta}_2$.

F-Test of Nested Models in R

If we want an F-test between two models, with the first model being the null hypothesis M_0 , we can use the `anova()` command (model 1 is just the previous model without `Total_Oil_Income_PC`

```
anova(model1, model2)
```

Analysis of Variance Table

Model 1: polity_2 ~ Total_Oil_Income_PC

Model 2: polity_2 ~ GDP_Per_Cap_Haber_Men_2 + Total_Oil_Income_PC

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	10268	537568				
2	10267	452505	1	85064	1930	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The first row of the output refers to the null hypothesis model.

The second row of the output refers to the alternate hypothesis. We can see that this row has the F -statistic and the p-value.

- We can see that model 2 is statistically significant (consult interpretation from the previous sections).

Chapter 5

Other Explanatory Variables

5.1 Polynomial Transformations

Sometimes, a linear (straight-line) best-fit line is a poor description of a relationship. For example, the relationship between two variables could be curved, not straight.

We can model more flexible relationships that are not straight lines, by including a transformation of the variable X that we are interested in. The most common transformations for non-linearity are polynomial transformations, including quadratic and cubic transformations.

Quadratic Transformations

Quadratic transformations of X take the following form:

$$E[Y] = \alpha + \beta_1 X + \beta_2 X^2$$

If you recall from high-school algebra, an equation that takes the form of $y = ax^2 + bx + c$ creates a *parabola*. Indeed, this transformation fits a parabola as the best-fit line. However, there are a few things to consider:

- A true parabola has a domain of $(-\infty, \infty)$. However, our model often does not need to do this. The best-fit parabola is only used for the range of plausible X values, given the nature of our explanatory variable. For example, if X was age, a negative number would make no sense.
- Because the parabola's domain often exceeds our plausible range of X values, the vertex of the parabola (where it changes directions) may not be in our data.
- We always include lower degree terms in our model. For example, in this quadratic (power 2) model, we also include the X term without the square.

To fit a model like this, we simply do the same process of minimising the sum of squared errors.

When we run a quadratic model in R, we get two coefficients: β_1 is attached to the X term, while β_2 is attached to the X^2 term. How do we interpret these coefficients?

- β_1 's value is no longer directly interpretable. This is because we cannot “hold all other coefficients constant”, since β_2 also contains the same X variable. Thus, we cannot isolate the effect of X and β_1 .
- β_2 's value also cannot be directly interpreted. However, β_2 can tell us two things. First, if the coefficient of β_2 is statistically significant, we can conclude that there is a non-linear relationship between X and Y . Second, if β_2 is negative, the best-fit parabola will open downwards, and if β_2 is positive, the best-fit parabola will open upwards.

If we want to interpret the magnitude of the model, we are best off using predicted values of Y (obtained using the model equation above).

There is one more thing we can interpret with the quadratic transformation: the **vertex** of the best-fit parabola. The vertex, if we remember our algebra, is either the maximum or minimum point of a parabola. Thus, if we remember from calculus and optimisation, we can find the maximum and minimums through setting the derivative equal to 0. For the quadratic model, this is as follows - we first find the derivative, then set the derivative equal to 0:

$$\begin{aligned}\hat{Y} &= \hat{\alpha} + \hat{\beta}_1 X + \hat{\beta}_2 X^2 \\ \frac{d\hat{Y}}{dX} &= 0 + \hat{\beta}_1 + 2\hat{\beta}_2 X \\ 0 &= \hat{\beta}_1 + 2\hat{\beta}_2 X \\ -\hat{\beta}_1 &= 2\hat{\beta}_2 X \\ X &= \frac{-\hat{\beta}_1}{2\hat{\beta}_2}\end{aligned}$$

This point is useful, as it is either the maximum or minimum of our best-fit parabola. This means that at the X value we calculate from this equation, we will either see the highest or lowest expected Y value.

General Polynomial Models

While quadratic models are the most common polynomial transformation, we do not have to stop there. We can continue to add further polynomials (although anything beyond cubic is exceedingly rare):

- Cubic: $E[Y] = \alpha + \beta_1 X + \beta_2 X^2 + \beta_3 X^3$
- Quartic: $E[Y] = \alpha + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 X^4$

Each higher order coefficient, if statistically significant, indicates that the relationship between X and Y , is not of the previous highest power.

- For example, if the cubic term β_3 is statistically significant, we can reject a quadratic relationship between X and Y

Remember to always include the lower power monomials within our polynomial model. For example, if you have a quartic transformation, you must also have the linear, quadratic, and cubic terms.

5.2 Logarithmic Transformations

Logarithmic transformations are another form of non-linear transformations. Logarithmic transformations are commonly used for heavily skewed variables, such as when the explanatory variable is income, wealth, and so on.

In situations with heavily skewed variables, we often replace X in our models with $\log(X)$. Note that in statistics, when we refer to logarithms, we are referring to natural logarithms, such that $\log(X) = \ln(X)$.

Thus, the logarithmic transformation takes the following form:

$$E[Y] = \alpha + \beta \log(X)$$

Interpretation of the β coefficient can be a little bit trickier for logarithmic transformations. We could interpret it in the same way we interpret linear regressions: given a one unit increase in the log of X , there is an expected β change in Y .

However, this issue is that this does not really say much - I mean, who knows what a *one unit increase in the log of X* even means?

However, with some properties of logarithms, we can actually create a more useful interpretation. Based on logarithm rules, we know the following to be true:

$$\begin{aligned}\log(X) + A &= \log(X) + \log(e^A) \\ &= \log(e^A \times X)\end{aligned}$$

Now, let us plug this into our original regression model:

$$\begin{aligned}E[Y|X] &= \alpha + \beta \log(X) \\ E[Y|e^A \times X] &= \alpha + \beta \log(e^A \times X) \\ &= \alpha + \beta [\log(X) + A] \\ &= \alpha + \beta A + \beta \log(X)\end{aligned}$$

Now find the difference between $E[Y|e^A \times X]$ and $E[Y|X]$:

$$\begin{aligned}E[Y|e^A \times X] - E[Y|X] &= [\alpha + \beta A + \beta \log(X)] - [\alpha + \beta \log(X)] \\ E[Y|e^A \times X] - E[Y|X] &= \beta A\end{aligned}$$

Thus, we can see that when we multiply X by e^A , we get an expected βA change in Y .

We can make this interpretation more useful by purposely choosing some value A that makes e^A make more sense. For example, if $A = 0.095$, then $e^A = 1.1$. Why is that A value useful? Well, that means when we multiply $X \times e^A$, we are actually doing $X \times 1.1$, which if you remember your percentages, means a 10% increase in X . Thus, increasing X by 10% is associated with an expected change of 0.095β units of Y .

5.3 Binary Explanatory Variable

Binary explanatory variables will change the interpretations of our coefficients. We can “solve” for these interpretations given the standard linear model $E[Y] = \alpha + \beta X$, given X has two categories $X = 0, X = 1$:

$$\begin{aligned}E[Y|X = 0] &= \alpha + \beta(0) = \alpha \\E[Y|X = 1] &= \alpha + \beta(1) = \alpha + \beta \\E[Y|X = 1] - E[Y|X = 0] &= (\alpha + \beta) - \alpha = \beta\end{aligned}$$

From the above, we can see the following interpretations of our coefficients:

- α is the expected value of Y given an observation in category $X = 0$
- $\alpha + \beta$ is the expected value of Y given an observation in category $X = 1$
- β is the expected difference in Y between the categories $X = 1$ and $X = 0$

Thus, we can see that β is measuring the difference between the two categories. In fact, β actually becomes a difference-in-means test, meaning that if β is statistically significant, we can conclude a significant difference in the mean Y between the two categories.

Because of the unique coefficient meanings in a regression with binary explanatory variables, the OLS estimator also takes a shortcut when estimated the coefficients α and β :

- Coefficient α is estimated as $\hat{\alpha} = \bar{Y}_0$
- Coefficient β is estimated as $\hat{\beta} = \bar{Y}_1 - \bar{Y}_0$

Where \bar{Y}_1 is the sample mean of Y for observations in category $X = 1$, and \bar{Y}_0 is the sample mean of Y for observations in category $X = 0$.

5.4 Polytomous Explanatory Variable

A **polytomous** variable is one with 3 or more categories that are unranked. A classic example is the variable *country*, which is a categorical variable with all the different countries included in a dataset such as Argentina, France, Mexico, etc.

How do we run a regression with polytomous explanatory variables? What happens is that we divide the variables into a set of dummy binary variables.

- Dummy binary variables are created for all except one of the categories in our variable. Each dummy variable has two values - 1 meaning the observation is in the category, and 0 meaning the observation is not in that category.
- The category without a dummy variable is the **reference/baseline** category. Essentially, when all other dummy variables are equal to 0, that is referring to the reference/baseline category.

Thus, for the n number of categories in X , we would create $n - 1$ dummy variables, and input it into a regression equation as follows:

$$E[Y] = \alpha + \beta_{x=0}X_{x=0} + \beta_{x=1}X_{x=1} + \dots + \beta_{x=n-1}X_{x=n-1}$$

For example, take the following polytomous variable: *company*, which contains the categories *microsoft*, *google*, and *apple*.

- Let us create dummy variables for 2 of the 3 categories.
- *Google* will become the first dummy variable X_g . When $X_g = 1$, that observation is part of the *google* category. When $X_g = 0$, that observation is NOT a part of the *google* category.
- *Apple* will become the second dummy variable X_a . When $X_a = 1$, that observation is part of the *apple* category. When $X_a = 0$, that observation is NOT a part of the *apple* category.
- *Microsoft* will not get its own dummy variable. This is because when both *apple* and *microsoft* $X_g = X_a = 0$ that is referring to the *microsoft* category (since these are the only observations not a part of either previous category).

Mathematically, this is how it would be represented in a regression equation:

$$E[Y] = \alpha + \beta_g X_g + \beta_a X_a$$

To find the expected value of each category, we would do the following:

$$\begin{aligned} E[Y|X = \text{Google}] &= E[Y|X_g = 1, X_a = 0] = \alpha + \beta_g(1) + \beta_a(0) = \alpha + \beta_g \\ E[Y|X = \text{Apple}] &= E[Y|X_g = 0, X_a = 1] = \alpha + \beta_g(0) + \beta_a(1) = \alpha + \beta_a \\ E[Y|X = \text{Microsoft}] &= E[Y|X_g = 0, X_a = 0] = \alpha + \beta_g(0) + \beta_a(0) = \alpha \end{aligned}$$

Thus, from these above equations, we can see the interpretation of the coefficients:

- α is the expected value of the reference category, in this case, *microsoft*.
- β_g is the expected Y difference between the *google* category and the reference category *microsoft*. The statistical significance of this coefficient would be a difference of means test between the two categories.
- β_a is the expected Y difference between the *apple* category and the reference category *microsoft*. The statistical significance of this coefficient would be a difference of means test between the two categories.

More generally, the β_j of category j 's dummy variable, represents the expected difference in Y between category j and the reference/baseline category.

Notice how the coefficient p -values are a difference-of-means test between two categories, and not a statistical significance test of the entire categorical variable *company* that has 3 different categories. To test the statistical significance of the entire categorical variable, we use an F-test.

An F-test, simply, tests a model's explanatory power against the explanatory power of a model where all coefficients $\alpha, \beta_1, \dots, \beta_n$ all are equal to 0. Thus, an F-test would allow us to test the significance of the effect of the variable *company* on Y , which our individual β coefficients do not tell us.

5.5 Interaction Effects

Interactions, also called moderating effects, means that the effect of some X_j on Y is not constant, and depends on some third variable X_k . Essentially, X_k 's value changes the relationship between X_j and Y .

This is quite common in the real world. For example, imagine outcome variable Y to be the severity of a car crash. X_1 can represent the darkness of the road at the time of the car crash. X_2 can represent the slipperiness of the road at the time of the car crash. We could quite reasonably expect that as the road is more slippery, i.e. X_2 increases, the darkness of the road X_1 's effect on the severity of a car crash Y might be stronger, since slippery further enhances the danger of dark roads.

Or for a more political example, Y could be the chance of a civil war occurring, X_1 is the severity of an economic crash, and X_2 is the development level of a country. We could quite reasonably expect that in the effect of a economic crash on a chance of civil war would be significantly higher in developing nations rather than developed. Or in other words, the chance that a civil war occurs due to a economic crash is higher in countries like Venezuela, North Korea and Eritrea, compared to the relationship in Norway, Switzerland, and Denmark.

Interaction effects are represented by two variables being multiplied together in a regression equation. In the model below, X_1 and X_2 are interacting with each other:

$$E[Y] = \alpha + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2$$

We can mathematically show that the effect of X_1 on Y is not constant - and varies due to the value of X_2 . We show this through finding the partial derivative of X_1 on Y , since the derivative is, by definition, the function of the rate of change between X_1 and Y .

$$\begin{aligned}\hat{Y} &= \hat{\alpha} + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 + \hat{\beta}_3 X_1 X_2 \\ \frac{\partial \hat{Y}}{\partial X_1} &= 0 + \hat{\beta}_1 + 0 + \hat{\beta}_3 X_2 \\ \frac{\partial \hat{Y}}{\partial X_1} &= \hat{\beta}_1 + \hat{\beta}_3 X_2\end{aligned}$$

As you can see, the relationship between X_1 and Y here depends on the value of X_2 . In more intuitive words, given a one unit increase in X_1 , there is an expected $\hat{\beta}_1 + \hat{\beta}_3 X_2$ increase in Y .

We can also find the effect of X_2 on Y :

$$\begin{aligned}\hat{Y} &= \hat{\alpha} + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 + \hat{\beta}_3 X_1 X_2 \\ \frac{\partial \hat{Y}}{\partial X_2} &= \hat{\beta}_2 + \hat{\beta}_3 X_1\end{aligned}$$

With these equations, we can interpret the coefficients of our model:

- $\hat{\beta}_1$ is the relationship between X_1 and Y , given $X_2 = 0$.
- $\hat{\beta}_2$ is the relationship between X_2 and Y , given $X_1 = 0$.

- $\hat{\beta}_3$ represents two things. For every one unit increase of X_2 , the magnitude of the relationship between X_1 and Y changes by $\hat{\beta}_3$. Similarly, for every one unit increase of X_1 , the magnitude of the relationship between X_2 and Y changes by $\hat{\beta}_3$
- α is still the expected value of Y when all explanatory variables equal 0.

The coefficient β_3 's significance level tells us if there is a statistically significant interaction. If β_3 is not statistically significant, we can often remove the interaction term. However, if β_3 is statistically significant, that means we have found two terms that interact.

Often times, our moderating effect X_2 is a binary variable (for example, developed/developing country, true/false, yes/no). In this scenario:

- $\hat{\beta}_1$ is the relationship between X_1 and Y when X_2 is in the category $X = 0$.
- $\hat{\beta}_1 + \beta_3$ is the relationship between X_1 and Y when X_2 is in the category $X = 1$.
- $\hat{\beta}_3$ is the difference in the magnitude of the relationship between X_1 and Y , between the categories $X = 1$ and $X = 0$.

You can get all of these interpretations above simply by plugging in $X_2 = 0$ and $X_2 = 1$ into the previous equations we have found.

5.6 Explanatory Variables in R

Part III

Regression for Causal Inference

Chapter 6

Gauss-Markov Assumptions

Chapter 7

Panel and Clustered Data

7.1 Hierarchical Data

7.2 Fixed Effects

7.3 Further Approaches

Chapter 8

Matching Regression

Chapter 9

Weighted Regression

Chapter 10

Partial Identification and Sensitivity Analysis

Part IV

Instrumental Variables Estimator

Chapter 11

Instrumental Variables Framework

Chapter 12

Other Instrumental Variable Designs

12.1 Continuous Independent Variables

12.2 Examiner Designs

12.3 Shift-Share Bartik Instruments

12.4 Recentered Instruments

Part V

Quasi-Experimental Methods

Chapter 13

Regression Discontinuity

Chapter 14

Simple Differences-in-Differences

Chapter 15

Generalised Differences-in-Differences

Chapter 16

Survey Experiments